

PREDICTIVE ANALYSIS WITH DECISION TREES

Submitted by:

Name : Sahana S

Register Number : 727723EUCS200

Class : CSE – ‘C’

Year : III

Department : Computer Science and Engineering

1. Introduction

Predictive analysis involves analyzing historical data to make predictions about future outcomes. Decision Trees are popular machine learning models used for predictive analytics due to their simplicity, interpretability, and ability to handle both classification and regression problems.

2. Objectives

- Implement Decision Tree for classification
- Understand entropy and Gini index
- Apply pruning techniques to avoid overfitting
- Visualize the decision tree
- Evaluate feature importance

3. Dataset Description

The Heart Disease dataset consists of patient medical records used to predict the presence of heart disease. The dataset contains numerical and categorical attributes related to heart health.

4. Decision Tree Theory

4.1 Entropy

Entropy measures the level of uncertainty or impurity in a dataset. Lower entropy indicates purer data.

$$Entropy(S) = -\sum p_i \log_2(p_i)$$

Where p_i is the probability of class i .

4.2 Information Gain

Used to decide the best feature to split.

$$IG(S, A) = Entropy(S) - \sum \frac{|S_v|}{|S|} Entropy(S_v)$$

4.3 Gini Index

Measures impurity (used by CART algorithm).

$$Gini = 1 - \sum p_i^2$$

Lower Gini → better split.

5. Overfitting in Decision Trees

Decision trees tend to memorize training data leading to overfitting.

Solutions:

- **Pre-Pruning:** Stop tree growth early
- **Post-Pruning:** Remove unnecessary branches

6. Methodology

1. Load dataset
2. Preprocess data
3. Compute entropy and Gini manually
4. Train Decision Tree
5. Apply pruning
6. Evaluate performance

7. Visualize tree
8. Interpret feature importance

7. Tools & Technologies

- Python
- NumPy
- Pandas
- Scikit-learn
- Matplotlib
- Graphviz

8. Results

- Accuracy before pruning
- Accuracy after pruning
- Improved generalization
- Key features identified

9. Conclusion

Decision Trees are powerful and interpretable machine learning models. By using entropy, Gini index, and pruning techniques, model performance can be optimized and overfitting can be avoided.

10. Future Enhancements

- Use Random Forests
- Apply Gradient Boosting
- Hyperparameter tuning

Source Code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# 1. IMPORT LIBRARIES
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder

# 2. LOAD DATASET
df = pd.read_csv("heart.csv")

print("\nDataset Preview:")
display(df.head())
print("\nDataset Shape:", df.shape)

# FIX: HANDLE CATEGORICAL DATA
```

```
label_encoder = LabelEncoder()

for col in df.columns:
    if df[col].dtype == 'object':
        df[col] = label_encoder.fit_transform(df[col])
```

```
print("\nDataset After Encoding:")
display(df.head())
```

3. MATHEMATICAL CALCULATIONS

```
def entropy(y):
    values, counts = np.unique(y, return_counts=True)
    probabilities = counts / counts.sum()
    return -np.sum(probabilities * np.log2(probabilities))
```

```
def gini_index(y):
    values, counts = np.unique(y, return_counts=True)
    probabilities = counts / counts.sum()
    return 1 - np.sum(probabilities ** 2)
```

```
print("\nEntropy of Target:", entropy(df['target']))
print("Gini Index of Target:", gini_index(df['target']))
```

4. FEATURE & TARGET SPLIT

```
X = df.drop('target', axis=1)
y = df['target']
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42  
)
```

5. DECISION TREE (NO PRUNING)

```
dt_full = DecisionTreeClassifier(  
    criterion='gini',  
    random_state=42  
)
```

```
dt_full.fit(X_train, y_train)  
y_pred_full = dt_full.predict(X_test)  
  
acc_full = accuracy_score(y_test, y_pred_full)  
print("\nAccuracy (No Pruning):", acc_full)
```

6. PRE-PRUNING

```
dt_pre_pruned = DecisionTreeClassifier(  
    criterion='gini',  
    max_depth=4,  
    min_samples_split=10,  
    random_state=42  
)
```

```
dt_pre_pruned.fit(X_train, y_train)  
y_pred_pre = dt_pre_pruned.predict(X_test)  
  
acc_pre = accuracy_score(y_test, y_pred_pre)  
print("Accuracy (Pre-Pruning):", acc_pre)
```

```
# 7. POST-PRUNING

path = dt_full.cost_complexity_pruning_path(X_train, y_train)
ccp_alphas = path ccp_alphas

dt_post_pruned = DecisionTreeClassifier(
    random_state=42,
    ccp_alpha=ccp_alphas[5]
)

dt_post_pruned.fit(X_train, y_train)
y_pred_post = dt_post_pruned.predict(X_test)

acc_post = accuracy_score(y_test, y_pred_post)
print("Accuracy (Post-Pruning):", acc_post)

# 8. CLASSIFICATION REPORT

print("\nClassification Report (Pre-Pruned Model):")
print(classification_report(y_test, y_pred_pre))

# 9. DECISION TREE VISUALIZATION

plt.figure(figsize=(20, 10))
plot_tree(
    dt_pre_pruned,
    feature_names=X.columns,
    class_names=['No Disease', 'Disease'],
    filled=True
)
plt.title("Decision Tree Visualization (Pre-Pruned)")
```

```
plt.show()

# 10. FEATURE IMPORTANCE

feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': dt_pre_pruned.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importance:")
display(feature_importance)

# 11. FEATURE IMPORTANCE VISUALIZATION

plt.figure(figsize=(10, 5))

plt.bar(feature_importance['Feature'], feature_importance['Importance'])

plt.xticks(rotation=90)
plt.title("Feature Importance Based on Decision Tree")
plt.tight_layout()
plt.show()

# 12. ACCURACY COMPARISON

models = ['No Pruning', 'Pre-Pruning', 'Post-Pruning']
accuracies = [acc_full, acc_pre, acc_post]

plt.figure(figsize=(6, 4))

plt.bar(models, accuracies)
plt.ylabel("Accuracy")
plt.title("Decision Tree Model Accuracy Comparison")
plt.show()

print("\nPROJECT EXECUTED SUCCESSFULLY")
```

Screenshots:

The screenshot shows the Jupyter Notebook interface with the following details:

- Left Sidebar (Explorer):** Shows the "OPEN EDITORS" section with "Untitled17.ipynb" selected.
- Top Bar:** Displays the title "jupiter", tabs for "Untitled17.ipynb" and "Welcome", and a Python version indicator "Python 3.14.0".
- Code Editor:** Contains the following Python code:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# 1. IMPORT LIBRARIES
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder

# 2. LOAD DATASET
df = pd.read_csv("heart.csv")

print("\nDataset Preview:")
display(df.head())
print("\nDataset Shape:", df.shape)

# FIX: HANDLE CATEGORICAL DATA
label_encoder = LabelEncoder()

for col in df.columns:
    if df[col].dtype == 'object':
        df[col] = label_encoder.fit_transform(df[col])

print("\nDataset After Encoding:")
display(df.head())

# 3. MATHEMATICAL CALCULATIONS

def entropy(y):
    values, counts = np.unique(y, return_counts=True)
    probabilities = counts / counts.sum()
    return -np.sum(probabilities * np.log2(probabilities))

def gini_index(y):
    values, counts = np.unique(y, return_counts=True)
```

At the bottom of the code editor, there are buttons for "Spaces: 4 LF Cell 3 of 3 Go Live Prettier" and a "Run" button.

The screenshot shows the Jupyter Notebook interface with the following details:

- Left Sidebar (Explorer):** Shows the "OPEN EDITORS" section with "Untitled17.ipynb" selected.
- Top Bar:** Displays the title "jupiter", tabs for "Untitled17.ipynb" and "Untitled17.ipynb > import numpy as np", and a Python version indicator "Python 3.14.0".
- Code Editor:** Contains the following Python code:

```
def entropy(y):
    values, counts = np.unique(y, return_counts=True)
    probabilities = counts / counts.sum()
    return -np.sum(probabilities * np.log2(probabilities))

def gini_index(y):
    values, counts = np.unique(y, return_counts=True)
    probabilities = counts / counts.sum()
    return 1 - np.sum(probabilities ** 2)

print("\nEntropy of Target:", entropy(df['target']))
print("Gini Index of Target:", gini_index(df['target']))

# 4. FEATURE & TARGET SPLIT
X = df.drop('target', axis=1)
y = df['target']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# 5. DECISION TREE (NO PRUNING)
dt_full = DecisionTreeClassifier(
    criterion='gini',
    random_state=42
)

dt_full.fit(X_train, y_train)
y_pred_full = dt_full.predict(X_test)

acc_full = accuracy_score(y_test, y_pred_full)
print("\nAccuracy (No Pruning):", acc_full)

# 6. PRE-PRUNING
dt_pre_pruned = DecisionTreeClassifier(
    criterion='gini',
    max_depth=4,
    min_samples_split=10,
    random_state=42
)
```

At the bottom of the code editor, there are buttons for "Spaces: 4 LF Cell 3 of 3 Go Live Prettier" and a "Run" button.

The screenshot shows a Jupyter Notebook interface with the following details:

- Left Sidebar (Explorer):** Shows the file structure with "Untitled17.ipynb" selected.
- Top Bar:** Includes tabs for "jupiter", "Untitled17.ipynb", and "Welcome".
- Code Editor:** Displays Python code for a decision tree project, including:
 - Utility functions: `gini_index` and `entropy`.
 - Data splitting: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42).
 - Decision Tree creation: `dt_full = DecisionTreeClassifier(criterion='gini', random_state=42)`.
 - Fitting and prediction: `dt_full.fit(X_train, y_train)` and `y_pred_full = dt_full.predict(X_test)`.
 - Accuracy calculation: `acc_full = accuracy_score(y_test, y_pred_full)`.
 - Pruning setup: `dt_pre_pruned = DecisionTreeClassifier(criterion='gini', max_depth=4, min_samples_split=10, random_state=42)`.
- Bottom Bar:** Shows "Spaces: 4 LF Cell 3 of 3 Go Live Prettier" and a "Python" icon.

The screenshot shows a Jupyter Notebook interface with the following details:

- Left Sidebar (Explorer):** Shows the file structure with "Untitled17.ipynb" selected.
- Top Bar:** Includes tabs for "jupiter", "Untitled17.ipynb", and "Welcome".
- Code Editor:** Displays Python code for a decision tree project, including:
 - Visualization setup: `plt.title("Decision Tree Visualization (Pre-Pruned)")` and `plt.show()`.
 - Feature importance: `feature_importance = pd.DataFrame({'Feature': X.columns, 'Importance': dt_pre_pruned.feature_importances_}).sort_values(by='Importance', ascending=False)`.
 - Print statement: `print("\nFeature Importance:")` followed by `display(feature_importance)`.
 - Feature importance visualization: `plt.figure(figsize=(10, 5))`, `plt.bar(feature_importance['Feature'], feature_importance['Importance'])`, `plt.xticks(rotation=90)`, `plt.title("Feature Importance Based on Decision Tree")`, `plt.tight_layout()`, and `plt.show()`.
 - Accuracy comparison: `models = ['No Pruning', 'Pre-Pruning', 'Post-Pruning']` and `accuracies = [acc_full, acc_pre, acc_post]`.
 - Comparison visualization: `plt.figure(figsize=(6, 4))`, `plt.bar(models, accuracies)`, `plt.ylabel("Accuracy")`, `plt.title("Decision Tree Model Accuracy Comparison")`, and `plt.show()`.
 - Success message: `print("\nPROJECT EXECUTED SUCCESSFULLY")`.
- Bottom Bar:** Shows "Spaces: 4 LF Cell 3 of 3 Go Live Prettier" and a "Python" icon.

EXPLORER

OPEN EDITORS

Untitled17.ipynb (Selected)

Welcome

Code **Markdown** **Run All** **Restart** **Clear All Outputs** **Jupyter Variables** **Outline** **...**

Python 3.14.0

Dataset Preview:

```
...    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  ca  thal  target
...  0   52   1   0     125   212   0     1    168     0     1.0     2   2   3   0
  1   53   1   0     140   203   1     0    155     1     3.1     0   0   3   0
  2   70   1   0     145   174   0     1    125     1     2.6     0   0   3   0
  3   61   1   0     148   203   0     1    161     0     0.0     2   1   3   0
  4   62   0   0     138   294   1     1    106     0     1.9     1   3   2   0
...
... Dataset Shape: (1025, 14)
... Dataset After Encoding:
```

```
...    age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  slope  ca  thal  target
...  0   52   1   0     125   212   0     1    168     0     1.0     2   2   3   0
  1   53   1   0     140   203   1     0    155     1     3.1     0   0   3   0
  2   70   1   0     145   174   0     1    125     1     2.6     0   0   3   0
  3   61   1   0     148   203   0     1    161     0     0.0     2   1   3   0
  4   62   0   0     138   294   1     1    106     0     1.9     1   3   2   0
...
... Entropy of Target: 0.9994994187527655
... Gini Index of Target: 0.49965306365258777
...
... Accuracy (N Pruning): 0.9853658536585366
... Accuracy (Pre-Pruning): 0.8
... Accuracy (Post-Pruning): 0.9414634146341463
...
... Classification Report (Pre-Pruned Model):
      precision    recall   f1-score   support
...
        0       0.88      0.70      0.78     102
        1       0.75      0.90      0.82     103
...
accuracy          0.81      0.80      0.80     205
macro avg       0.81      0.80      0.80     205
weighted avg    0.81      0.80      0.80     205
```

Spaces: 4 LF Cell 3 of 3 Go Live Prettier

EXPLORER

OPEN EDITORS

Untitled17.ipynb (Selected)

Welcome

Code **Markdown** **Run All** **Restart** **Clear All Outputs** **Jupyter Variables** **Outline** **...**

Python 3.14.0

Decision Tree Visualization (Pre-Pruned)

Feature Importance:

Feature	Importance
2 cp	0.407309
12 thal	0.195828
11 ca	0.158818
9 oldpeak	0.100518
10 slope	0.031802
7 thalach	0.029911

Spaces: 4 LF Cell 3 of 3 Go Live Prettier

